

Tooling in the age of serverless computing

Donna Malayeri
Product and Community Manager, Pulumi
[@lindydonna](https://twitter.com/lindydonna)

In 2008, AWS had 4 services

Products

Infrastructure Services

- › **Amazon Elastic Compute Cloud (Amazon EC2)**
Amazon Elastic Compute Cloud delivers scalable, pay-as-you-go compute capacity in the cloud.
- › **Amazon SimpleDB**
Amazon SimpleDB works in conjunction with Amazon S3 and Amazon EC2 to run queries on structured data in real time.
- › **Amazon Simple Storage Service (Amazon S3)**
Amazon Simple Storage Service provides a fully redundant data storage infrastructure for storing and retrieving any amount of data, at any time, from anywhere on the Web.
- › **Amazon Simple Queue Service (Amazon SQS)**
Amazon Simple Queue Service provides a hosted queue for storing messages as they travel between computers, making it easy to build automated workflow between Web services.



Virtual Machines

Database

Storage

Queues

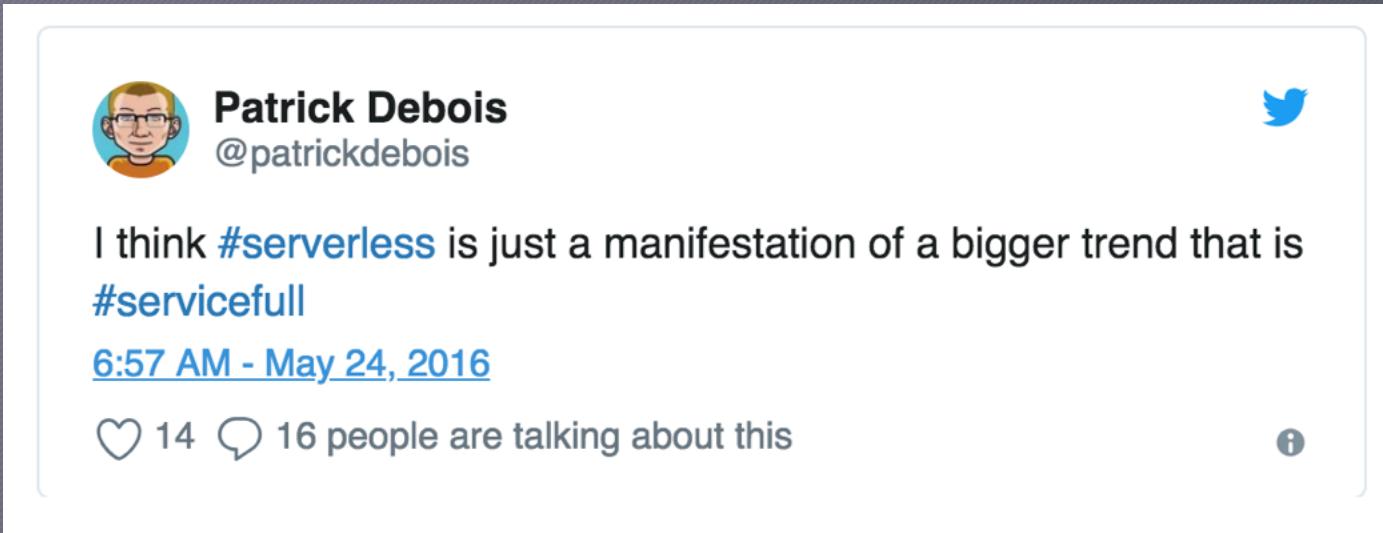
Today, each vendor has over a hundred services

The image displays three screenshots of cloud provider websites side-by-side, illustrating the complexity of modern cloud offerings:

- AWS (left):** Shows the "Explore Our Products" page with a sidebar for "Compute" services like EC2, Lambda, and Lambda.
- Azure (center):** Shows the "Products" section with categories like "Featured", "Virtual Machines", "Azure SQL Database", "Azure Cosmos DB", "Azure Kubernetes Service (AKS)", and "Cognitive Services".
- Google Cloud Platform (right):** Shows a comprehensive list of services under various categories such as Compute, Storage & Databases, Cloud AI, and Big Data, including services like Cloud Storage, Cloud Functions, Cloud Spanner, and Cloud ML Engine.

Each screenshot includes a yellow header bar with the provider's name. The Azure screenshot also features a yellow bar at the top with the word "AZURE".

#serverless? Or #servicefull?



A screenshot of a Twitter post from Patrick Debois (@patrickdebois). The post contains the following text: "I think #serverless is just a manifestation of a bigger trend that is #servicefull". It includes a timestamp "6:57 AM - May 24, 2016", engagement metrics "14", and "16 people are talking about this". The Twitter logo is visible in the top right corner of the card.

Patrick Debois
@patrickdebois

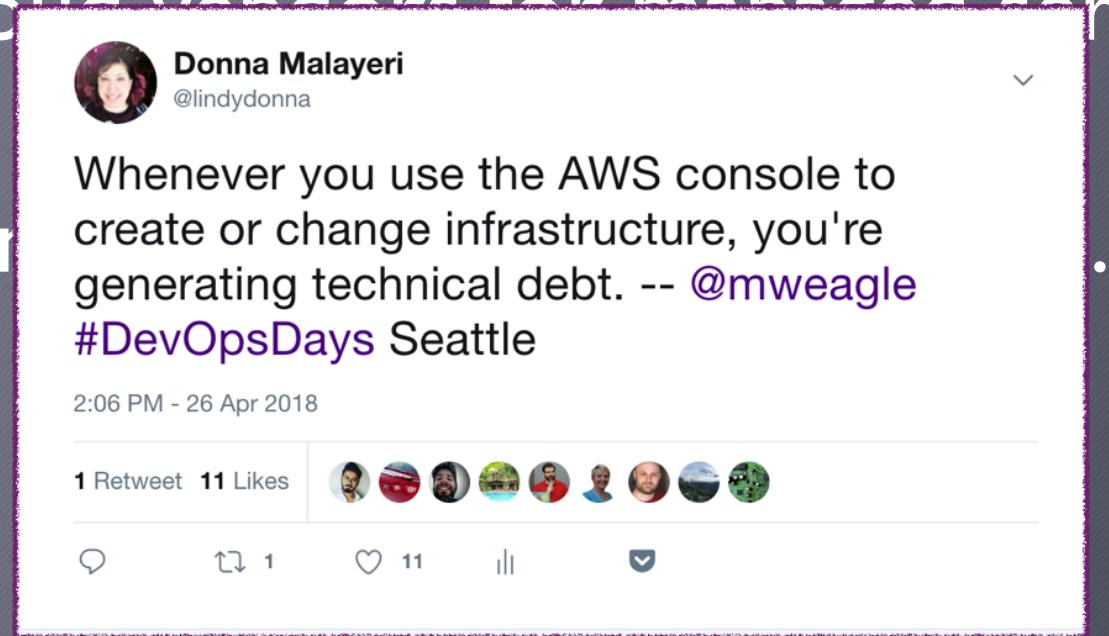
I think **#serverless** is just a manifestation of a bigger trend that is **#servicefull**

[6:57 AM - May 24, 2016](#)

14 16 people are talking about this

Now, we have a new problem

- Relying on cloud vendors for managed services == net good
 - Having to manage



@lindydonna

Option A: use vendor's technology

```
Policies:
- PolicyDocument:
  Statement:
  - Action:
    - logs:CreateLogGroup
    - logs:CreateLogStream
    - logs:PutLogEvents
    Effect: Allow
    Resource: "*"
    Sid: AllowLogging
    Version: '2012-10-17'
  PolicyName: LambdaWriteCWLogs
- PolicyDocument:
  Statement:
  - Action:
    - s3:Get*
    Effect: Allow
    Resource:
      Fn::Sub: arn:aws:s3:::${PhotoRepoS3Bucket}/*
    Sid: ReadFromPhotoRepoS3Bucket
    Version: '2012-10-17'
  PolicyName: ReadFromPhotoRepoS3Bucket
- PolicyDocument:
  Statement:
  - Action:
    - s3:PutObject
    Effect: Allow
    Resource:
      Fn::Sub: arn:aws:s3:::${PhotoRepoS3Bucket}/*
    Sid: WriteToPhotoRepoS3Bucket
    Version: '2012-10-17'
  PolicyName: WriteToPhotoRepoS3Bucket
- PolicyDocument:
  Statement:
  - Action:
    - dynamodb:UpdateItem
    - dynamodb:PutItem
    Effect: Allow
    Resource:
      Fn::Sub: arn:aws:dynamodb:${AWS:Region}:${AWS:AccountId}:table/${ImageMetadataDBTable}
    Sid: WriteToImageMetadataDBTable
    Version: '2012-10-17'
  PolicyName: WriteToImageMetadataDBTable
- PolicyDocument:
  Statement:
  - Action:
    - rekognition:DetectLabels
    Effect: Allow
    Resource: "*"
    Sid: RekognitionDetectLabels
    Version: '2012-10-17'
  PolicyName: RekognitionDetectLabels
- PolicyDocument:
  Statement:
  - Action:
    - states:StartExecution
    Effect: Allow
    Resource: "*"
    Sid: StepFunctionStartExecution
    Version: '2012-10-17'
  PolicyName: StepFunctionStartExecution
Type: AWS::IAM::Role
CreatesEventTriggerFunction:
Properties:
  CodeUrl: s3://media-sharing-refarch/8c7e2179b5bc3489407509396c78b95e
Description: Used with CloudFormation as a custom resource helper to enable S3 event trigger to invoke the start step function Lambda function.
```

AWS CLOUDFORMATION

```
"variables": {
  "storageAccountName": "[concat(uniquestring(resourceGroup().id), '-sacustmdate')]",
  "imagePublisher": "Canonical",
  "imageOffer": "UbuntuServer",
  "nicName": "networkInterface",
  "vnName": "vni1",
  "virtualNetworkName": "virtualNetwork1",
  "publicIPaddressName": "publicip1",
  "addressPrefix": "10.0.0.0/16",
  "subnet1": "Subnet-1",
  "subnetPrefix": "10.0.0.0/24",
  "vmStorageAccountContainerName": "vhds",
  "publicIPaddressType": "Dynamic",
  "storageAccountType": "Standard_LRS",
  "vnetID": "[resourceId('Microsoft.Network/virtualNetworks', variables('virtualNetworkName'))]",
  "subnetRef": "[concat(variables('vnetID'), '/subnets/', variables('subnet1Name'))]",
  "apiVersion": "2015-06-15"
},
"resources": [
  {
    "type": "Microsoft.Storage/storageAccounts",
    "name": "[variables('storageAccountName')]",
    "apiVersion": "2015-06-15",
    "location": "[resourceGroup().location]",
    "properties": {
      "accountType": "[variables('storageAccountType')]"
    }
  },
  {
    "apiVersion": "2015-06-15",
    "type": "Microsoft.Network/publicIPAddresses",
    "name": "[variables('publicIPaddressName')]",
    "location": "[resourceGroup().location]",
    "properties": {
      "publicIPAllocationMethod": "[variables('publicIPAddressType')]",
      "dnsSettings": {
        "domainNameLabel": "[parameters('dnsLabelPrefix')]"
      }
    }
  },
  {
    "apiVersion": "2015-06-15",
    "type": "Microsoft.Network/virtualNetworks",
    "name": "[variables('virtualNetworkName')]",
    "location": "[resourceGroup().location]",
    "properties": {
      "addressSpace": {
        "addressPrefixes": [
          "[variables('addressPrefix')]"
        ]
      },
      "subnets": [
        {
          "name": "[variables('subnet1Name')]",
          "properties": {
            "addressPrefix": "[variables('subnet1Prefix')]"
          }
        }
      ]
    }
  }
],
"outputs": {
  "apiVersion": "2015-06-15",
  "type": "Microsoft.Network/networkInterfaces",
  "name": "[variables('nicName')]"
},
```

AZURE RESOURCE MANAGER

```
imports:
  - path: path/to/my_vm_template.jinja
    name: my_renamed_template.jinja
  - path: special_vm.py
```

If your template uses other templates as dependencies, import the dependent templates in your configuration as well:

```
imports:
  - path: path/to/my_vm_template.jinja
  - path: special_vm.py
  - path: base_vm.jinja
```

You can also import text files in order to inline the content. For example, if you create a file named `resource_type.txt` with the following string:

```
compute.v1.instance
```

Import it into your configuration and provide the content inline like so:

```
imports:
  - path: resource_type.txt

resources:
  - name: my-vm
    type: {{ imports["resource_type.txt"] }} # Resolves to "compute.v1.instance"
    properties:
      zone: us-central1-a
      machineType: zones/us-central1-a/machineTypes/f1-micro
      disks:
        - deviceName: boot
          type: PERSISTENT
          boot: true
          autoDelete: true
          initializeParams:
            sourceImage: projects/debian-cloud/global/images/family/debian-8
      networkInterfaces:
        - network: global/networks/default
          accessConfigs:
            - name: External NAT
              type: ONE_TO_ONE_NAT
```

GOOGLE CLOUD DEPLOYMENT MANAGER

Option B: use a deployment tool



When should you use a tool?

- Your team does not have expertise in the vendor technology
 - Or, some team members don't have this expertise
- You're comfortable taking a new dependency, if it buys productivity
- You have a small team and can't afford to specialize

A framework for evaluating tools

How do I
describe
infrastructure?

How do I
write my app?

How do I
deploy?

How do I
learn?

Serverless deployment tools



Option B: use a deployment tool



Option B: use a deployment tool



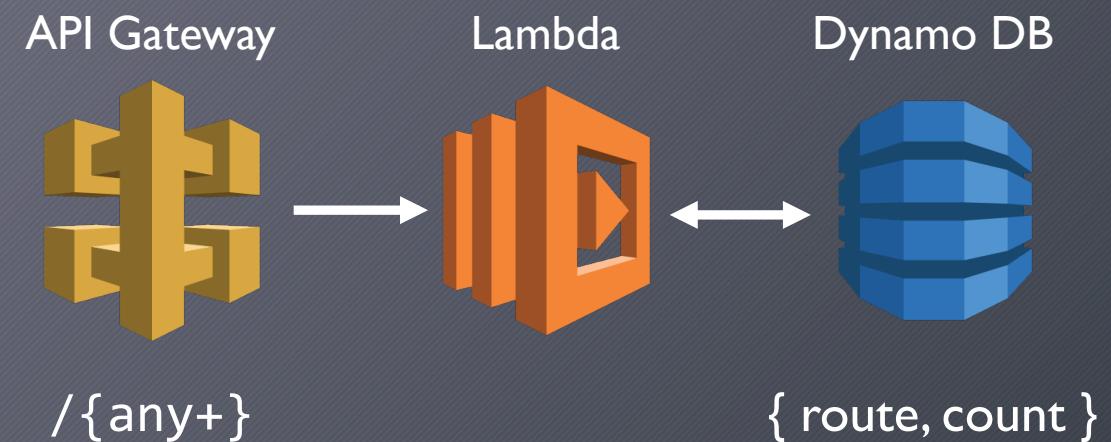
More at github.com/lindydonna/velocity-examples

Count the number of times a route has been hit

```
$ curl $URL/foo  
{"id": "foo", "count": 1}
```

```
$ curl $URL/foo  
{"id": "foo", "count": 2}
```

```
$ curl $URL/bar  
{"id": "bar", "count": 1}
```





Serverless Framework

- *Serverless is your toolkit for deploying and operating serverless architectures. Focus on your application, not your infrastructure.*
- First class experience for serverless functions on AWS, Azure, GCP, and more
- Manages code and infrastructure
- Can author functions in JavaScript, Python, Java, and more
- Rich ecosystem of community plugins

Serverless Framework example

```
provider:  
  name: dynamoDbTable  
  region: us-east-1  
  type: dynamoDBTable  
  deletionPolicy: delete  
  properties:  
    attributeDefinitions:  
      - attributeName: id  
        attributeType: S  
    keySchema:  
      - attributeName: id  
        keyType: H  
    provisionedThroughput:  
      readCapacityUnits: 1  
      writeCapacityUnits: 1  
    tableName: ${self:provider.environment.DYNAMO_DB_TABLE}
```

45 LOC for config

```
const AWS = require('aws-sdk');

const dynamoDb = new AWS.DynamoDB.DocumentClient();

module.exports.handler = (event, context, callback) => {
  ...
}
```

30 LOC for Lambda

Deploying with Serverless Framework

- `sls deploy` - Uses vendor's deployment service
 - Such as CloudFormation on AWS

After deployment

- `sls logs` - View logs for a specific function
- CloudFormation console

The screenshot shows the Serverless Platform Beta dashboard. At the top, there is a navigation bar with links to Dashboard, Service routes-api, Stage dev, and Function app. Below the navigation, a user profile for 'lindydonna / routes-api' is shown, with options for 'Private' and 'Share'. A dropdown menu indicates the stage is 'dev' (Last deployed 2 days ago). The main content area displays a Lambda function named 'app' (nodejs6.10), which was last deployed 2 days ago. Below this, detailed configuration information is provided: ARN (arn:aws:lambda:us-east-1:153052954103:function:routes-api-dev-app), Memory (1024), Timeout (6), and Endpoints (ANY - https://idn52tqt21.execute-api.us-east-1.amazonaws.com and ANY - https://idn52tqt21.execute-api.us-east-1.amazonaws.co). The final section shows the function documentation status: 'This function hasn't been documented yet.'

Serverless Framework

How do I
describe
infrastructure?

How do I
write my app?

How do I
deploy?

How do I
learn?

- Lambda and APIG are first-class
- For everything else, use CloudFormation

- No change to how Lambda is authored

- `sls deploy`, uses CloudFormation

- An active community
- Lots of plugins and tutorials

Terraform

- *HashiCorp Terraform enables you to safely and predictably create, change, and improve infrastructure. Define infrastructure as code to increase operator productivity and transparency.*
- Defines a custom language for describing deployments to AWS, Azure, etc
- No special support for serverless scenarios

Terraform example

```
resource "aws_lambda_function" "example" {
    name        = "ServerlessExample"
    runtime     = "nodejs8.10"
    handler     = "app.handler"
    role        = "${aws_iam_role.lambda_exec.arn}"
    source_code_hash = "${base64sha256(file("${data.archive_file.lambda.output_path}"))}"
    publish     = true
    environment {
        variables = {
            DYNAMODB_TABLE = "${aws_dynamodb_table.dynamo-table.id}"
        }
    }
}
```

150 LOC for config

30 LOC for Lambda

Deploying with Terraform

`terraform apply`

- Uses vendor APIs directly, does not use services like CloudFormation
 - Has a checkpoint file to store state of last deployment
 - In free version, you manage this yourself

After deployment

- No specific tooling for viewing resources or getting logs

Terraform

How do I
describe
infrastructure?

How do I
write my app?

How do I
deploy?

How do I
learn?

- Use HCL to configure any cloud resource on any cloud
- Can author your own provider

- No change to how Lambda is authored

- Uses vendor APIs directly
- You have to manage a state file

- Lots of tutorials and examples for infrastructure
- Less common for serverless



Claudia.js

- *Claudia makes it easy to deploy Node.js projects to AWS Lambda and API Gateway. It automates all the error-prone deployment and configuration tasks, and sets everything up the way JavaScript developers expect out of the box.*
- Supports JavaScript on AWS only
- Automatically sets up CORS, request/response format, etc
- With API builder, use API Gateway with regular JavaScript web server syntax
- Manages only API Gateway and Lambda resources

Claudia.js example

Create Dynamo table manually

```
aws dynamodb create-table --table-name RouteCounts ...
```

```
{
  "DYNAMODB_TABLE"
}
```

```
claudia create
--api-module index --policies policy \
--set-env-from-json env.json
```

policy.json

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
```

18 LOC for config

Deploying with Claudia.js

- Deploy new app with `claudia create`
 - Uses AWS APIs directly, creates a JSON file
- Deploy code with `claudia update`

After deployment

- No specific tooling for viewing resources or getting logs

Claudia.js

How do I
describe
infrastructure?

How do I
write my app?

How do I
deploy?

How do I
learn?

- JSON for APIG and Lambda
- Anything else, use an additional tool

- Convenient syntax for defining APIs

- Uses AWS APIs directly
- You have to manage a json state file

- Lots of tutorials and examples
- Upcoming book later this year



Create containers, serverless functions, and cloud infrastructure, with all the benefits of immutable infrastructure, and real programming languages.

- For cloud configuration, use JavaScript, Python or Go
- For functions and containers, use any supported language
- Manage any resource in AWS, Azure, Kubernetes, or GCP
- Use low-level resources or higher-level framework

Pulumi example: low level

```
const aws = require("@pulumi/aws")

const policy = { ... };

const deployment = new aws.apigateway.Deployment("api-deployment", {
    restApi: restApi,
    stageName: "",
});

const stage = new aws.apigateway.Stage("api-stage", {
    restApi: restApi,
    deployment: deployment,
    stageName: stageName
});

const invokePermission = new aws.lambda.Permission("api-lambda-permission", {
    action: "lambda:invokeFunction",
    function: lambda,
    principal: "apigateway.amazonaws.com",
    sourceArn: deployment.executionArn.apply(arn => arn + "*/*"),
});
```

110 LOC for config

30 LOC for Lambda

Pulumi example: using @pulumi/aws-serverless

```
const aws = require("@pulumi/aws");
const policy = { ... };
const role = new aws.iam.Role(...);
const fullAccess = new aws.iam.PolicyAttachment(...);

let lambda = codegen(role, handler, runtimes);
const api = new serverless.apigateway.API("myapi", {
    routes: [
        { method: "GET", path: "/{proxy+}", handler: lambda },
    ],
});
environment: {
    variables: {
        "DYNAMODB_TABLE": table.name
    }
},
});
```

54 LOC for config

30 LOC for Lambda

Pulumi example: define Lambda *inline*

```
// The handler for GET /{route+}.
const getHandler = async (event) => {
  const awssdk = require("aws-sdk");
  const dynamo = new awssdk.DynamoDB.DocumentClient();

  const tableName = table.name.get();
  ...
}

const api = new serverless.apigateway.API("api", {
  routes: [
    { method: "GET", path: "/{route+}", handler: getHandler }
  ]
});
```

52 LOC for config
AND Lambda

Reference table object directly

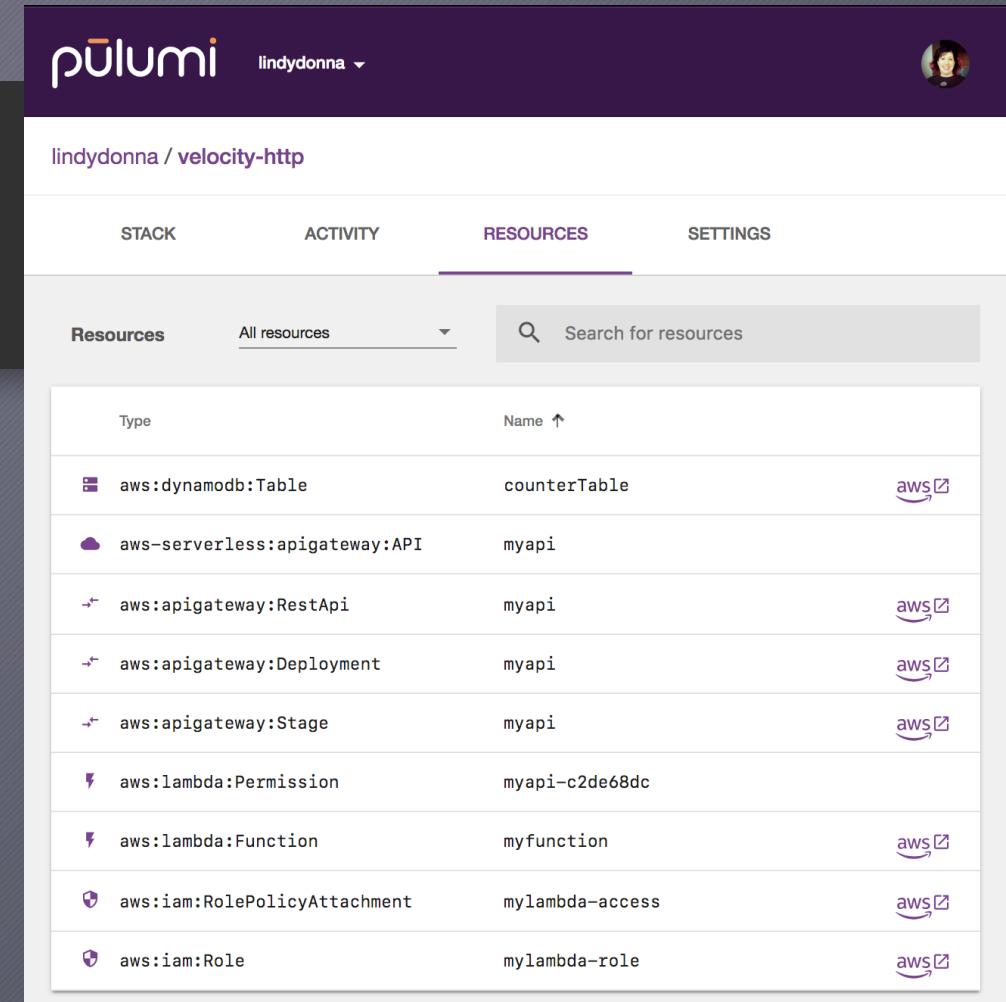
Deploying with Pulumi

pulumi update

- Uses vendor APIs directly
- Deployment state is stored on pulumi.com

After deployment
pulumi logs

```
$ pulumi logs -f
2018-05-25T12:57:30.705-07:00[          ] Running thumbnailer task.
2018-05-25T12:58:34.960-07:00[          ] Starting ffmpeg task...
```



The screenshot shows the Pulumi web interface with a dark purple header. The header includes the Pulumi logo, the user name "lindydonna", and a profile picture. Below the header, the page title is "lindydonna / velocity-http". There are four tabs: "STACK", "ACTIVITY", "RESOURCES" (which is underlined in purple), and "SETTINGS". Under the "RESOURCES" tab, there is a search bar with the placeholder "Search for resources". The main content area displays a table of resources:

| Type | Name | aws |
|-------------------------------|-----------------|---------------------|
| aws:dynamodb:Table | counterTable | aws |
| aws-serverless:apigateway:API | myapi | aws |
| aws:apigateway:RestApi | myapi | aws |
| aws:apigateway:Deployment | myapi | aws |
| aws:apigateway:Stage | myapi | aws |
| aws:lambda:Permission | myapi-c2de68dc | |
| aws:lambda:Function | myfunction | aws |
| aws:iam:RolePolicyAttachment | mylambda-access | aws |
| aws:iam:Role | mylambda-role | aws |

Pulumi

How do I
describe
infrastructure?

How do I
write my app?

How do I
deploy?

How do I
learn?

- Use JavaScript, Python or Go
- First class support for containers and serverless
- Target AWS, Azure, GCP, k8s

- Can author as regular Lambda
- Or inline with infra code

- Uses vendor APIs directly
- State file is saved on pulumi.com

- Still in beta, but lots of examples and tutorials

Summary

- For a more productive cloud experience, consider a deployment tool
- Answer these questions
 - *How do I describe **infrastructure**?*
 - *How do I write my **app**?*
 - *How do I **deploy**?*
 - *How do I **learn**?*
- All examples available at github.com/lindydonna/velocity-examples
- Try out Pulumi at www.pulumi.com